

## Selected Topics in AI: Week 6

### Autoencoders

- **Autoencoders** were first introduced with the goal of learning to **reconstruct** the input observations  $x_i$  with the **lowest error possible**.
- Why would one want to learn to reconstruct the input observations?
  - An autoencoder would be an algorithm that can give as output an image that is as similar as possible to the input one.

**Definition 1** *An autoencoder is a type of algorithm with the primary purpose of learning an "informative" representation of the data that can be used for different applications<sup>a</sup> by learning to reconstruct a set of input observations well enough.*

- **Another Definition:** Autoencoders are a specific type of feedforward neural networks where **the input is the same as the output**. They **compress** the input into a **lower-dimensional code** and then **reconstruct the output** from this representation. The code is a compact “summary” or “compression” of the input, also called the latent-space representation.

- **Architecture of the Autoencoders**




Figure 1: General structure of an autoencoder.

- 
- The autoencoders' main components are three:
  - an encoder,
  - a latent feature representation,
  - and a decoder.
- Generally speaking, we want the autoencoder to **reconstruct the input well enough**.
- Still, at the same time, it should create a latent representation (the output of the encoder part in Figure 1) that is useful and meaningful.
- For example, latent features on hand-written digits could be the number of lines required to write each number or the angle of each line and how they connect.
- While learning, **we extract the essential information that will allow us to solve a problem (writing digits, for example)**.
- In most typical architectures, **the encoder and the decoder are neural networks**.
- In general, **the encoder can be written as a function  $g$**  that will depend on some parameters.

$$\mathbf{h}_i = g(\mathbf{x}_i)$$

- Where  $h_i \in \mathbb{R}^q$  (the latent feature representation) is the output of the encoder block in Figure 1 when we evaluate it on the input  $x_i$ .
- The decoder (and the output of the network that we will indicate with  $\tilde{x}_i$ ) can be written then as a second generic function  $f$  of the latent features.

$$\tilde{x}_i = f(h_i) = f(g(x_i))$$

- Training an autoencoder simply means finding the functions  $g(\cdot)$  and  $f(\cdot)$  that satisfy

$$\arg \min_{f,g} \langle [\Delta(x_i, f(g(x_i)))] \rangle$$

- 
- where  $\Delta$  indicates a measure of how the input and the output of the autoencoder differ (basically our loss function will penalize the difference between input and output) and  $\langle \cdot \rangle$  indicates the average over all observations.

## Deep Generative Models and Latent Variables

- **Deep Generative Modeling** is used to **model real-world data distribution**. This type of learning has many applications in synthesizing new data, anomaly detection, learning semantically rich representations of data, and more.
- There are various types of deep generative models, based on how they are trained, and how they learn the data distribution.
- Let's suppose we have a collection of human faces. For example, we have CelebA (CelebFaces Attributes) dataset. Below are some samples from the CelebA-HQ dataset.



- Now suppose looking at the above faces we want to make a new face. What will be the intuitive process? Well, we first draw a rough sketch for the structure of the face, then we decide on the facial attributes, hair color, etc, and finally, we will come up with a new face.
- So intuitively before making an image of a face, there are some **variables** in the image that are important before drawing a face.

- Such factors or variables are called **Latent Variables**.
- These variables **don't appear explicitly** but are **important in the data generation process**.
- So, when drawing a new face image, we go from a low-level representation of data (**Latent Variables**) to a high-level representation.
- Now, define the above scenario mathematically.
  - In that case, we have **high dimensional data**  $\mathbf{x} \in \mathbb{X}^D$  (face images, D-dimensional vector), and for each image, we have some **low dimensional latent variables**  $\mathbf{z} \in \mathbb{Z}^M$  (pose, face color, hairs, etc M-dimensional vector).
  - The generative process of a face image can be described as

$$\mathbf{z} \sim P(\mathbf{z}) \quad \mathbf{x} \sim P(\mathbf{x}|\mathbf{z}) \quad (1)$$

- The above notation denotes  $\mathbf{z}$  and  $\mathbf{x}$  are being sampled from respective probability distributions.
- What it means is that we first obtain some latent vector  $\mathbf{z}$  (for example, deciding facial attributes and all), and then we generate **a face image  $\mathbf{x}$  based on those latent vectors**.
- One thing to note is that both image and latent variables are sampled from **a probability distribution**.
- All these variables have a **certain valid range**, and there are certain values for each variable that are **more probable** than others. Hence, real-world variables can be thought of as **random variables following some probability distribution**.
- Similarly, in order to understand the **conditional distribution** above, let's say I pick a random height, for example, 174 cm, the probability of the individual being a male for that height would be higher than being a female as

generally males are taller. Hence, the occurrence of some variables affects the likelihood of other variables.

## Model and the Objective

- The idea of the latent variable model is that, if we have some distribution of latent variables  $z$  and we know the conditional distribution  $P(x|z)$ , we can get  $P(x)$  from the probability theory as follows:

$$P(x) = \int_z P(X, Z) dz = \int_z P(z) P(x|z) dz \quad (2)$$

- Now  $P(z)$  is prior.  $P(x|z)$  is a distribution whose parameters can be learned to **maximize the likelihood of the data under that distribution**.
- let's take prior ( $P(z)$ ) as **normal distribution  $N(0,1)$** , and  $P(x|z)$  a **Gaussian whose parameters (mean and sigma) are learned by neural networks** as a function of  $z$ .

$$P(x|z) = \mathcal{N}(\mu_\theta(z), \Sigma_\theta(z)) \quad (3)$$

- Hence, the marginal likelihood is now given as:

$$P(x) = \int_z P(z) P(x|z) dz = \int_z \mathcal{N}(0, 1) \mathcal{N}(\mu_\theta(z), \Sigma_\theta(z)) dz$$

- The objective is to maximize  **$P(x)$**  given the data  $x \in X^D$  ( $x$  is  $D$  dimensional). We have  $n$  data points (training points) that we assume are **independent** of each other, and the total likelihood of the dataset is given as the **product of the likelihood of each data point  $x^{(i)}$** .

- Normally, we maximize **log-likelihood**, which transforms the **product into summation** and the total log-likelihood is given as the **sum of the log-likelihood** of each data point  $x^{(i)}$ :

$$\max_{\theta} \log P_{\theta}(X) = \sum_{i=1}^n \log \int_z p_z(z) p_{\theta}(x^i|z) dz \quad (4)$$

- The  $\theta$  in the above equation comes from the fact, that we define  $P(x|z)$  as some distribution parameterized by  $\theta$ , which can be optimized to fit the distribution to the dataset.
- Let's first look at how to solve the above integral, once we have a solution for it, **we can derive the equation for maximizing  $P(x)$** .
- It can be inferred that the integral inside the log is not tractable (no analytical solution). Hence, we have to **numerically** integrate it. But normally, we work in high-dimensional spaces, so there is a **curse of dimensionality** and we cannot integrate it **numerically**.
- What else can we do? We can rewrite the above equation as **an expectation concerning  $P(z)$** .

$$\max_{\theta} \log P_{\theta}(X) = \sum_{i=1}^n \log \mathbb{E}_{z \sim P(z)} [p_{\theta}(x^i|z)] \quad (5)$$

•

## Latent Variable Models

- Generative models share characteristics. They are often unsupervised and don't make use of observed inputs.
- Generative models have multiple uses. They can estimate density or map between domains.
- There are at least three classes of generative models:
  - Autoregressive models
  - Generative adversarial models
  - **Latent variable models.**
- Latent variable models understand **what causes data**. They introduce an **unobserved latent variable  $z$**  and a **prior  $p(z)$** .
- They specify a likelihood  $p(x|z)$  from which observations can be generated —  $z \sim p(z)$ ,  $x \sim p(x|z)$ . This defines a joint distribution  $p(x,z) = p(x|z) * p(z)$ .